

CS492: Probabilistic Programming

Amortised Inference

Hongseok Yang
KAIST

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r)))
```

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

Importance sampling with proposal $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

Importance sampling with proposal $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$w_i = 1$

Importance sampling with proposal $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.


```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_i = 1$$

Importance sampling with **proposal** $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = 1 * p(.4)/q(.4)$$
$$r_1 = .4$$

Importance sampling with **proposal** $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

Importance sampling with proposal $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r)))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

Importance sampling with proposal $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

$$w_2 = .144 * p(.6)/q(.6)$$
$$r_2 = .6$$

Importance sampling with proposal $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

$$w_2 = .144 * p(.6)/q(.6)$$
$$r_2 = .6$$

Importance sampling with proposal $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

$$w_2 = .144 * p(.6)/q(.6)$$
$$r_2 = .6$$

Importance sampling with **proposal $q(r)$** .

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

How to find good q ?

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

$$w_2 = .144 * p(.6)/q(.6)$$
$$r_2 = .6$$

Importance sampling with proposal $q(r)$.

1. Generate $(w_1, r_1), \dots, (w_n, r_n)$ by running prog.

2. Estimate $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$.

How to find good q ? **Use amortised inference!**

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))]  
    a (observe (flip r) true)  
    b (observe (flip r) true)  
    c (observe (flip r) false)]  
    r))
```

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))]  
    a (observe (flip r) false)  
    b (observe (flip r) false)  
    c (observe (flip r) false)]  
    r))
```

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))]  
    a (observe (flip r) true)  
    b (observe (flip r) true)  
    c (observe (flip r) true)]  
    r))
```

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))]  
    a (observe (flip r) false)  
    b (observe (flip r) false)  
    c (observe (flip r) true)]  
    r))
```

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

Other examples:
Finance model,
captcha, brain, etc.

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

Other examples:
Finance model,
captcha, brain, etc.

Amortised inference.

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

Other examples:
Finance model,
captcha, brain, etc.

Amortised inference. 1) Learn a proposal $q(x; y)$ parameterized by obs. y via preprocessing.

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

Other examples:
Finance model,
captcha, brain, etc.

Amortised inference. 1) Learn a proposal $q(x; y)$ parameterized by obs. y via preprocessing. 2) Use $q(x; y_0)$ for any actual observation y_0 later.

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

Other examples:
Finance model,
captcha, brain, etc.

Amortised inference. 1) Learn a proposal $q(x; y)$ parameterized by obs. y via preprocessing. 2) Use $q(x; y_0)$ for any actual observation y_0 later.

neural nets

We often need to infer posterior of one model multiple times with different observations.

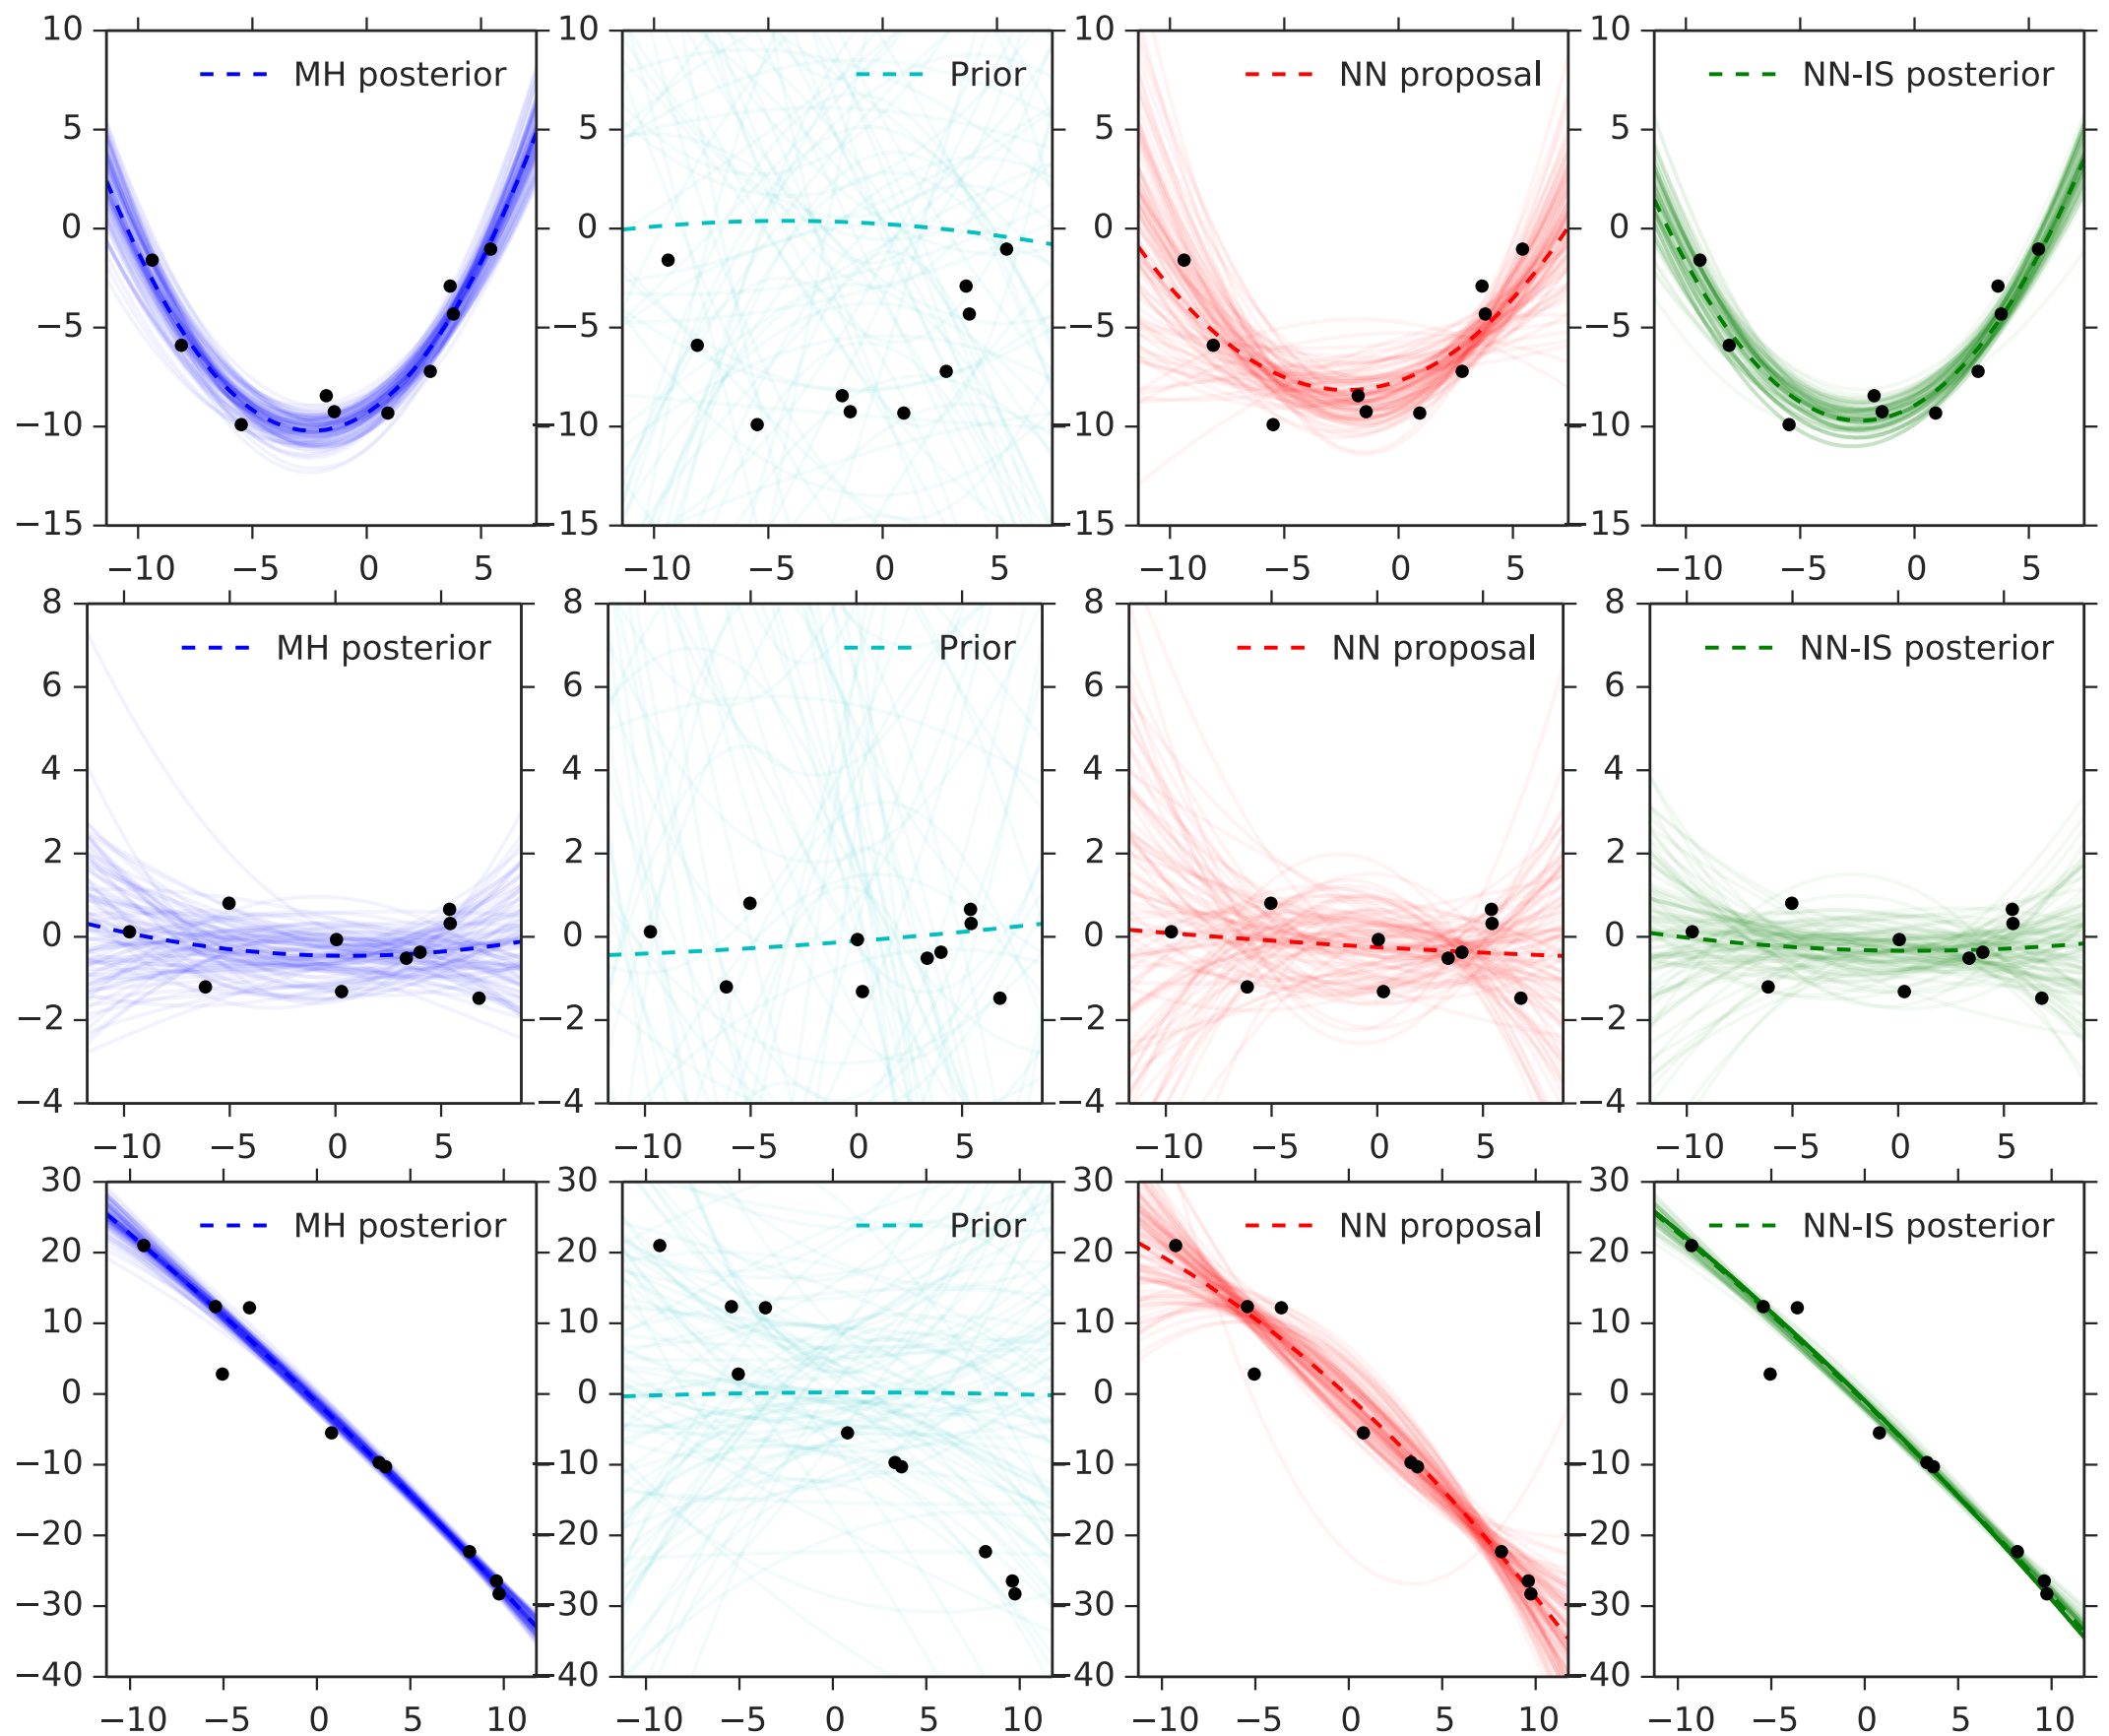
```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))]  
    a (observe (flip r) false)  
    b (observe (flip r) false)  
    c (observe (flip r) true)]
```

sample/object duality
and reverse KL

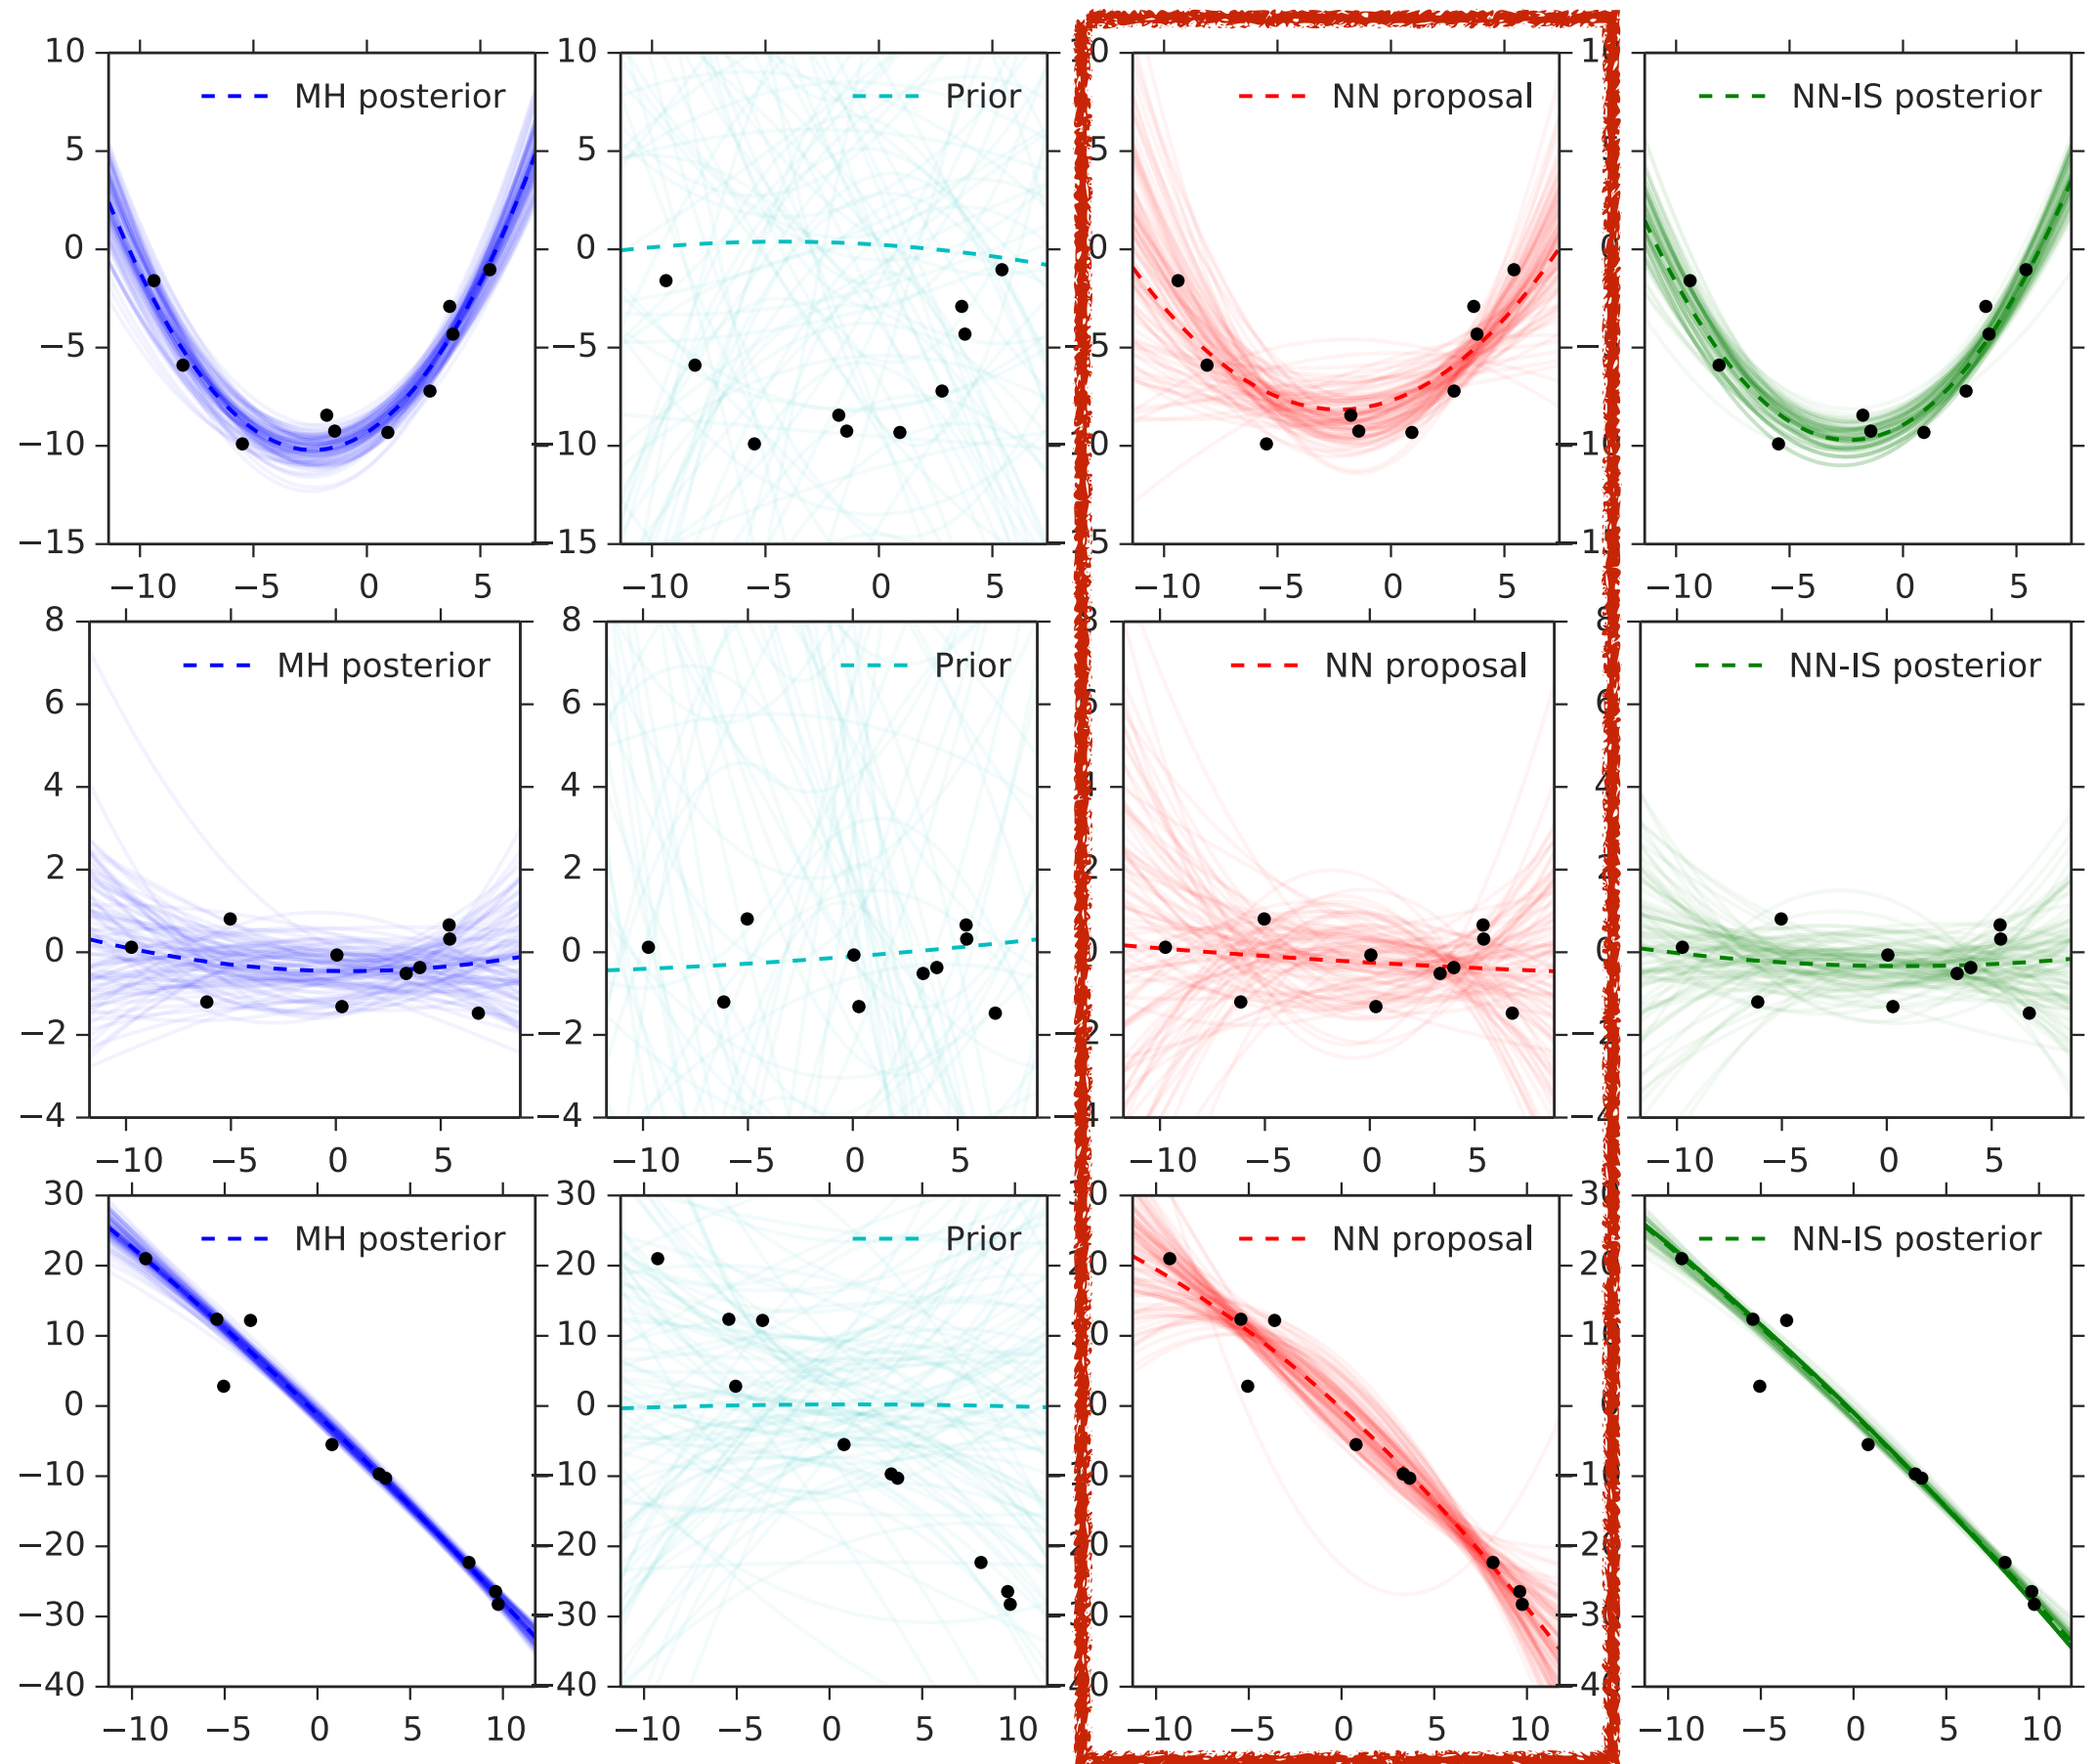
Other examples:
Finance model,
captcha, brain, etc.

Amortised inference. 1) **Learn** a proposal $q(x; y)$ parameterized by obs. y via preprocessing. 2) Use $q(x; y_0)$ for any actual observation y_0 later.

neural nets



Model for non-linear regression [Paige et al., ICML16]



Model for non-linear regression [Paige et al., ICML16]

Observed images


(W4kgvQ)


(uV7FeWB)


(MqhnpT)

more preprocessing

Samples



W4kgvQ

uV7EeWB

MqhnpT

WA4rjvQ

uV7FeWB

MypppT

Woxewd9

mTTEMMm

RIrpES

BKvu2Q

C9QDsoN

rS5FP2B

less preprocessing

Captcha solving [Le et al., AISTATS16]

Learning outcome

Can describe how amortised inference works for models written in math.

Can explain key ideas behind implementing amortised inference for probabilistic programs.

Proposal learning problem

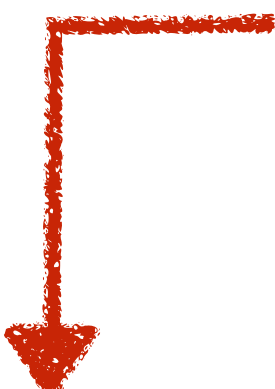
Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_\theta(x;y)$ parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y .

Proposal learning problem

Given:



Specified by $p(x)$ and $p(y|x)$.
Interested in $p(x|y)$.
But specific y not given yet.

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_\theta(x;y)$ parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y .

Proposal learning problem

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. **proposal $q_\theta(x;y)$** parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y .

Differentiable wrt. θ for fixed x,y .

E.g. $q_\theta(x;y) = \text{normal}(x; f_\theta(y), g_\theta(y))$ for neural nets f,g .

Proposal learning problem

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_\theta(x;y)$ parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y .

Proposal learning problem

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_\theta(x;y)$ parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y .

Proposal learning problem

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_\theta(x;y)$ parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y .

Proposal learning problem

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_\theta(x;y)$ parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y .

Proposal learning problem tackled by amortised inf.

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_\theta(x;y)$ parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y .

Proposal learning problem tackled by amortised inf.

Given:

y sampled
from $p(y)$

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_\theta(x;y)$ parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y . ←

Proposal learning problem tackled by amortised inf.

Given:

y sampled
from $p(y)$

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_\theta(x;y)$ parameterized by θ and y .

Find θ such that $q_\theta(x;y)$ is good for most y .

Small KL divergence from $p(x|y)$ to $q_\theta(x;y)$.

$$\text{KL}[p(x|y) \parallel q_\theta(x;y)] = \mathbb{E}_{p(x|y)}[\log(p(x|y)/q_\theta(x;y))].$$

Proposal learning problem

$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]]$.

Solve this by stochastic gradient descent.

inf.

y sampled
from $p(y)$

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_{\theta}(x;y)$ parameterized by θ and y .

Find θ such that $q_{\theta}(x;y)$ is good for most y .

Small KL divergence from $p(x|y)$ to $q_{\theta}(x;y)$.

$$\text{KL}[p(x|y) \parallel q_{\theta}(x;y)] = \mathbb{E}_{p(x|y)} [\log(p(x|y)/q_{\theta}(x;y))].$$

Proposal learning problem

$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]]$.

Solve this by stochastic gradient descent.

inf.

y sampled
from $p(y)$

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_{\theta}(x;y)$ parameterized by θ and y .

Find θ such that $q_{\theta}(x;y)$ is good for most y .

Small **KL divergence** from $p(x|y)$ to $q_{\theta}(x;y)$.

$$\text{KL}[p(x|y) \parallel q_{\theta}(x;y)] = \mathbb{E}_{p(x|y)} [\log(p(x|y)/q_{\theta}(x;y))].$$

KL divergence from $p_1(x)$ to $p_2(x)$

Denoted by $KL[p_1(x) \parallel p_2(x)]$.

$$KL[p_1(x) \parallel p_2(x)] := \mathbb{E}_{p_1(x)}[\log(p_1(x)/p_2(x))].$$

Average log ratio.

Measures how $p_1(x)$ is close to $p_2(x)$.

KL divergence from $p_1(x)$ to $p_2(x)$

Denoted by $KL[p_1(x) \parallel p_2(x)]$.

$$KL[p_1(x) \parallel p_2(x)] := \mathbb{E}_{p_1(x)}[\log(p_1(x)/p_2(x))].$$

Average log ratio.

Measures how $p_1(x)$ is close to $p_2(x)$.

[Q] What is $KL[p(x) \parallel p(x)]$?

KL divergence from $p_1(x)$ to $p_2(x)$

Denoted by $KL[p_1(x) \parallel p_2(x)]$.

$$KL[p_1(x) \parallel p_2(x)] := \mathbb{E}_{p_1(x)}[\log(p_1(x)/p_2(x))].$$

Average log ratio.

Measures how $p_1(x)$ is close to $p_2(x)$.

[Q] What is $KL[[t \mapsto .1; f \mapsto .9] \parallel [t \mapsto .5; f \mapsto .5]]$?

Proposal learning problem

$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]]$.

Solve this by stochastic gradient descent.

inf.

y sampled
from $p(y)$

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_{\theta}(x;y)$ parameterized by θ and y .

Find θ such that $q_{\theta}(x;y)$ is good for most y .

Small KL divergence from $p(x|y)$ to $q_{\theta}(x;y)$.

$$\text{KL}[p(x|y) \parallel q_{\theta}(x;y)] = \mathbb{E}_{p(x|y)} [\log(p(x|y)/q_{\theta}(x;y))].$$

Proposal learning problem

$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]]$.

Solve this by **stochastic gradient descent**.

inf.

y sampled
from $p(y)$

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_{\theta}(x;y)$ parameterized by θ and y .

Find θ such that $q_{\theta}(x;y)$ is good for most y .

Small KL divergence from $p(x|y)$ to $q_{\theta}(x;y)$.

$$\text{KL}[p(x|y) \parallel q_{\theta}(x;y)] = \mathbb{E}_{p(x|y)} [\log(p(x|y)/q_{\theta}(x;y))].$$

Stochastic gradient descent
for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Stochastic gradient descent
for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...

(until θ doesn't change)

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Learning rate

Initialise θ

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...

(until θ doesn't change)

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...



Can't compute, but can approximate.

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...



Can't compute, but can approximate.

Sample $(x_1, y_1), \dots, (x_n, y_n)$ from $p(x, y)$.

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...



Can't compute, but can approximate.

Sample $(x_1, y_1), \dots, (x_n, y_n)$ from $p(x, y)$.

$$\nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]] \approx -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i).$$

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...



Can't compute, but can approximate.

Sample $(x_1, y_1), \dots, (x_n, y_n)$ from $p(x, y)$.

$$\nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]] \approx -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i).$$

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...



Can't compute, but can approximate

Sample $(x_1, y_1), \dots, (x_n, y_n)$ from $p(x, y)$.

$$\nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]] \approx -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i).$$

exists since $q_{\theta}(x_i; y_i)$
is differentiable.



Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

Repeat the following until θ doesn't change:

1. Sample $(x_1, y_1), \dots, (x_n, y_n)$ from $p(x, y)$
2. $G \leftarrow -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i)$
3. $\theta \leftarrow \theta - 0.01 * G$

Proposal learning problem

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

Solve this by stochastic gradient descent.

inf.

y sampled
from $p(y)$

Given:

1. joint dist. $p(x,y)$ for latent x and observed y ,
2. proposal $q_{\theta}(x;y)$ parameterized by θ and y .

Find θ such that $q_{\theta}(x;y)$ is good for most y .

Small KL divergence from $p(x|y)$ to $q_{\theta}(x;y)$.

$$\text{KL}[p(x|y) \parallel q_{\theta}(x;y)] = \mathbb{E}_{p(x|y)} [\log(p(x|y)/q_{\theta}(x;y))].$$

What about probabilistic programs?

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

Repeat the following until θ doesn't change:

1. Sample $(x_1, y_1), \dots, (x_n, y_n)$ from $p(x, y)$
2. $G \leftarrow -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i)$
3. $\theta \leftarrow \theta - 0.01 * G$

Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise θ

Repeat the following until θ doesn't change:

1. Sample $(x_1, y_1), \dots, (x_n, y_n)$ from $p(x, y)$
2. $G \leftarrow -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i)$
3. $\theta \leftarrow \theta - 0.01 * G$

How to
sample y ?

Sample/observe duality

To sample observations, just replace sample by observe.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) true)]  
    r))
```

Sample/observe duality

To sample observations, just replace sample by observe.

```
(defquery biased-coin-joint []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (sample (flip r))  
        b (sample (flip r))  
        c (sample (flip r))]  
    [r [a b c]]))
```

Sample/observe duality

To sample observations, just replace sample by observe.

```
(defquery biased-coin-joint []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (sample (flip r))  
        b (sample (flip r))  
        c (sample (flip r))]  
    [r [a b c]]))
```

[Q] Define a sensible $q_{\theta}(r;a,b,c)$.

References

1. Inference networks for sequential Monte Carlo in graphical models. Paige et al. ICML'16.
2. Inference compilation and universal probabilistic programming. Let et al. AISTATS'17.